

ANKIT KAUSHIK

New delhi, Delhi | [Linkedin](#) | [Github](#) | [Portfolio](#) | [Leetcode](#) | +91 7011209041 | DeveloperAnkitKaushik@gmail.com

SUMMARY

Full-Stack & Gen AI Engineer with 2+ year building production LLM applications — AI agents, RAG pipelines, and workflow automation — across Python, Next.js, and Node. Focused on turning real business problems into **AI-assisted** products, from natural-language product discovery to multi-step agent workflows.

WORK EXPERIENCE

Software Engineer — [Wallmantra](#), New Delhi

May 2025 - Present

- Built and shipped "**Grace AI**" an AI-assisted shopping layer that lets customers find products in natural language and get personalized recommendations across the catalog.
- Added **LLM-powered customer support** workflows that handle common queries automatically and keep responses consistent, reducing manual support load.
- Improved storefront performance, cutting page load times by ~25% through caching and query optimization.
- Set up a **modular, reusable** service architecture with onboarding docs that helped new developers get productive faster.
- Worked with product, design, and growth teams to ship **AI features** that made products easier to discover and improved checkout.

Software Engineer Intern, [Amantya Technologies](#) (Onsite), Gurugram

March 2025 - May 2025

- Designed **multi-step agent workflows** that automated client request handling and reduced manual work for the operations team by ~40%.
- Built a structured logging and monitoring setup with automatic file rotation that significantly cut log storage and made issues easier to trace.
- Standardised log formats** and error-tracing across services, making debugging and deployments faster and more reliable for users

PROJECTS

Orange — Smart OpenAI-Compatible LLM Gateway & Cost Router

Python, Docker, FastAPI, OpenAI SDK, LLMs (Gemini/OpenAI), Embeddings, PyTorch/ONNX

- Smart Dynamic Routing & Cost Optimization:** Engineered an OpenAI-compatible gateway that analyzes incoming prompts to route simple tasks to lightweight models and complex tasks to heavy reasoning models, reducing LLM API costs by 19.5% while preserving 84% response quality.
- Ultra-Low Latency In-Process Classification:** Built a local embedding-based classifier with a 2-stage judge escalation fallback that detects query complexity in just ~16.2ms (P50) without adding any external network hops or sidecars.
- Zero-Refactor SDK Integration:** Architected the gateway as a drop-in proxy requiring only a single base_url swap in standard OpenAI SDKs (Python/TypeScript), maintaining 100% compatibility with streaming responses, usage tracking, and typed errors.
- Failover & Safe Caching Pipeline:** Implemented an exact-match caching layer with a 0.0% false-positive rate alongside automated provider circuit breakers and cost ceilings to guarantee high availability and budget protection.

Heads Up! — Early Product-Safety Intelligence Platform

Python, FastAPI, LangGraph, RAG, PostgreSQL/pgvector, Redis, Next.js, TypeScript, Docker, Crawl4AI

- Built a **multi-agent AI pipeline** (7 agents orchestrated with LangGraph) that investigates public web data — reviews, forums, marketplaces — to surface early product-safety risks before official recalls, producing a source-backed safety report.
- Engineered a real-time **streaming architecture** using FastAPI, Redis job queues, and Server-Sent Events so users watch each investigation stage (resolve → crawl → extract → cluster → score) update live in the Next.js dashboard.
- Implemented explainable, deterministic **risk scoring with RAG** (semantic retrieval over complaints via OpenAI/local embeddings + pgvector) and complaint clustering, keeping scores rule-based while the LLM only explains them — fully containerized with **Docker**.

EDUCATION

- Master of Computer Applications (MCA) | 92%**
Guru Gobind Singh Indraprastha University (GGSIPU), New Delhi
- Bachelor of Computer Applications (BCA) | 90%**
Guru Gobind Singh Indraprastha University (GGSIPU), New Delhi

September 2023 – May 2025

November 2020 - August 2023

ADDITIONAL INFORMATION

- Languages:** Python, **JavaScript (ES6+)**, C++, Java
- AI/LLM Development:** LangChain, RAG, Vector Databases (FAISS), Embeddings, Agent Workflows, LLM Evaluation, n8n
- Backend:** Node/Express, REST APIs, JWT Authentication, SQL
- Frontend:** React.js, **Next.js**, HTML5, CSS3
- Developer Tools & Platforms:** Git, Firebase, Postman, Figma
- Cloud & DevOps:** Docker, CI/CD (GitHub Actions), AWS, Git, Agile
- Certifications:** Microsoft Generative AI Engineering Specialisation, Namaste DSA, Namaste React, Namaste Node (Akshay Saini) & Meta Front-End Developer Specialization
- Awards/Activities:** Achieved **4899 World rank** out of 7 lakh participants in Naukri Young Turk, Former President of CodeShaala Tech Society (**Led 5 events**, mentored 30+ juniors), Participated in 3+ Hackathons, Led backend in all.
- Competitive Profile:** Solved **1000+ DSA** problems, **2150+ rating** on Leetcode with **139 world ranking** in a contest.